

Practical Machine Learning Project

Nicolas

Get and Clean Data

In this project, we are using data from the accelerometers on the belt, forearm, arm, and dumbbell of six participants to predict how well they did the exercise.

```
library(knitr)
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(rpart)
library(rpart.plot)
library(randomForest)
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
##
##     margin
```

```
library(corrplot)
```

Download the Data

```
trainUrl <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
testUrl <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"
trainCSV <- "./pml-training.csv"
testCSV <- "./pml-testing.csv"
if (!file.exists(trainCSV)) {
  download.file(trainUrl, destfile=trainCSV, method="curl")
}
if (!file.exists(testCSV)) {
  download.file(testUrl, destfile=testCSV, method="curl")
}
```

Read the Data

```
trainRaw <- read.csv("./pml-training.csv")
testRaw <- read.csv("./pml-testing.csv")
dim(trainRaw)
```

```
## [1] 19622 160
```

```
dim(testRaw)
```

```
## [1] 20 160
```

The training dataset contains 19622 observations and 160 variables, while the testing dataset contains 20 observations and 160 variables. The “classe” variable in the training dataset is the outcome to predict. **### Clean the data** Remove observations with missing values and non useful variables.

```
sum(complete.cases(trainRaw))
```

```
## [1] 406
```

Remove columns that contain NA missing values.

```
trainRaw <- trainRaw[, colSums(is.na(trainRaw)) == 0]
testRaw <- testRaw[, colSums(is.na(testRaw)) == 0]
```

Remove columns that do not contribute to the accelerometer measurements.

```
classe <- trainRaw$classe
trainRemove <- grepl("^X|timestamp|window", names(trainRaw))
trainRaw <- trainRaw[, !trainRemove]
trainCleaned <- trainRaw[, sapply(trainRaw, is.numeric)]
trainCleaned$classe <- classe
testRemove <- grepl("^X|timestamp|window", names(testRaw))
testRaw <- testRaw[, !testRemove]
testCleaned <- testRaw[, sapply(testRaw, is.numeric)]
```

Split the training data

Let's split the clean training dataset into training dataset at 70% and a validation dataset at 30%. The validation dataset will be used to conduct cross validation before the testing dataset.

```
set.seed(22519) # For reproducible purpose
inTrain <- createDataPartition(trainCleaned$classe, p=0.70, list=F)
trainData <- trainCleaned[inTrain, ]
valData <- trainCleaned[-inTrain, ]
```

Data Modeling

Fit a predictive model for the activity recognition using **Random Forest** algorithm as it automatically selects important variables and it is robust to correlated covariates and outliers.

```
controlRf <- trainControl(method="cv", 5)
modelRf <- train(classe ~ ., data=trainData, method="rf", trControl=controlRf,
ntree=250)
modelRf
```

```
## Random Forest
##
## 13737 samples
##    52 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10989, 10989, 10991, 10990, 10989
## Resampling results across tuning parameters:
##
##  mtry  Accuracy   Kappa
##    2    0.9901727 0.9875673
##   27    0.9917015 0.9895017
##   52    0.9840572 0.9798282
##
## Accuracy was used to select the optimal model using  the largest value.
## The final value used for the model was mtry = 27.
```

Estimate the performance of the model on the validation dataset.

```
predictRf <- predict(modelRf, valData)
confusionMatrix(valData$classe, predictRf)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1673    0    0    0    1
##           B    5 1131    3    0    0
##           C    0    0 1021    5    0
##           D    0    0   13  949    2
##           E    0    0    1    6 1075
##
## Overall Statistics
##
##           Accuracy : 0.9939
##           95% CI : (0.9915, 0.9957)
##           No Information Rate : 0.2851
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9923
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity           0.9970   1.0000   0.9836   0.9885   0.9972
## Specificity           0.9998   0.9983   0.9990   0.9970   0.9985
## Pos Pred Value        0.9994   0.9930   0.9951   0.9844   0.9935
## Neg Pred Value        0.9988   1.0000   0.9965   0.9978   0.9994
## Prevalence            0.2851   0.1922   0.1764   0.1631   0.1832
## Detection Rate        0.2843   0.1922   0.1735   0.1613   0.1827
## Detection Prevalence  0.2845   0.1935   0.1743   0.1638   0.1839
## Balanced Accuracy      0.9984   0.9992   0.9913   0.9927   0.9979
```

```
accuracy <- postResample(predictRf, valData$classe)
accuracy
```

```
## Accuracy      Kappa
## 0.9938828 0.9922620
```

```
err <- 1 - as.numeric(confusionMatrix(valData$classe, predictRf)$overall[1])
err
```

```
## [1] 0.006117247
```

The estimated accuracy of the model is 99.38% and the estimated out of sample error is 0.61%.

Predicting on Test Dataset

Apply the model to the original testing dataset.

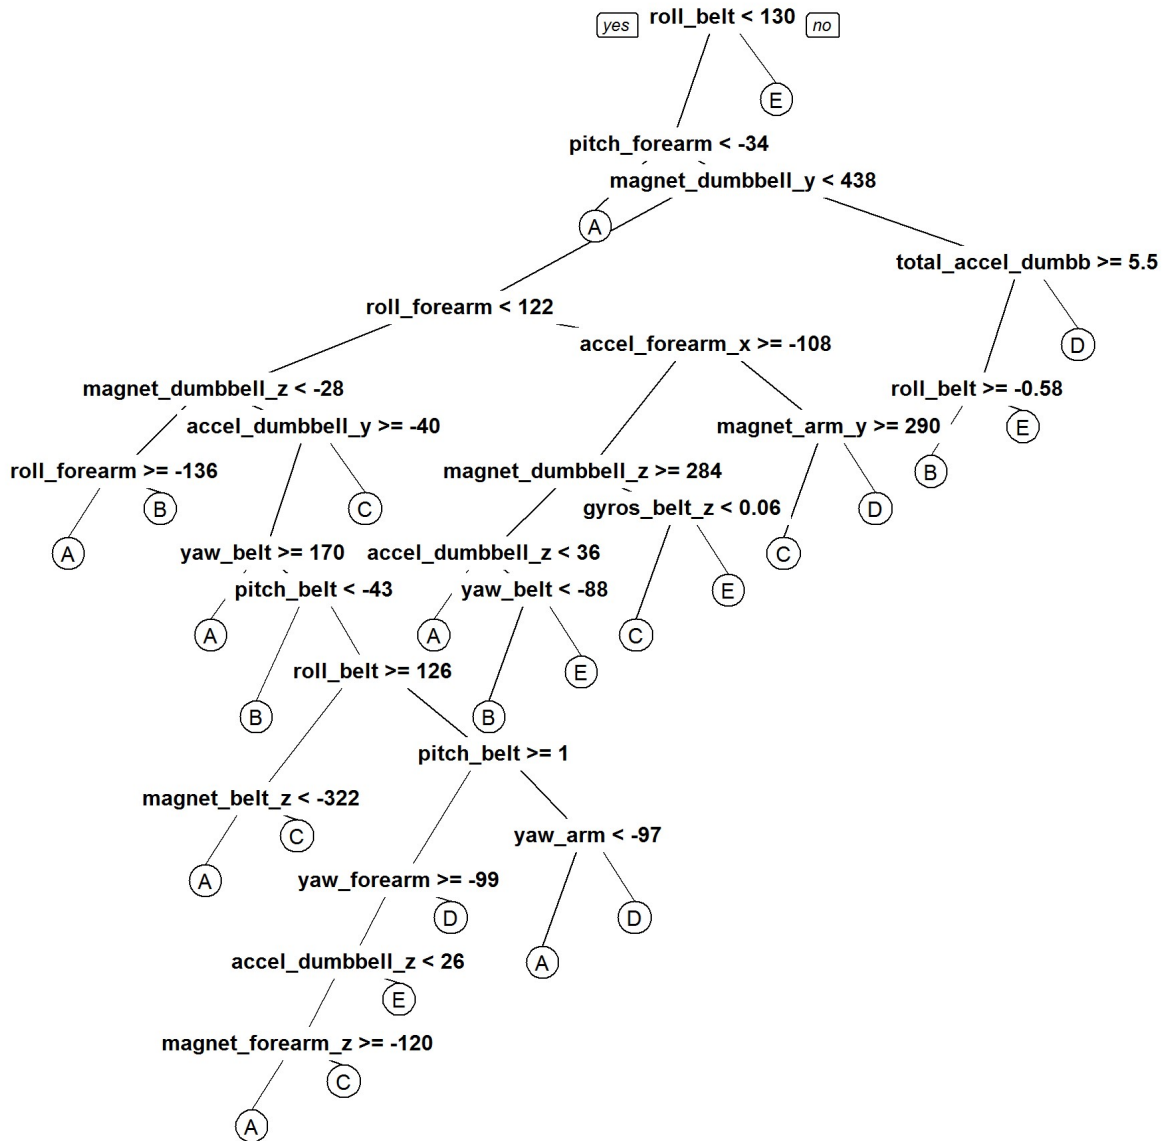
```
result <- predict(modelRf, testCleaned[, -length(names(testCleaned))])  
result
```

```
## [1] B A B A A E D B A A B C B A E E A B B B  
## Levels: A B C D E
```

Figures:

Decision Tree Visualization

```
treeModel <- rpart(classe ~ ., data=trainData, method="class")  
prp(treeModel)
```



Matrix Visualization of Correlation

```
corrPlot <- cor(trainData[, -length(names(trainData))])
corrplot(corrPlot, method="color")
```

